

Lecture 02: Tabular RL

Petr Kuderov

Researcher at MIPT & AIRI

February 04, 2026

Tabular Reinforcement Learning

- finite MDPs
- tabular state and action spaces
- episodic setting
- lecture + interleaved practice ... again
- data: interaction trajectories
- goal: *optimal policy via value estimates* (value-based RL family)

Prologue

- reversed order vs standard RL path
 - start from “final” algorithms
 - then derive where they come from
-
- first: what code does
 - later: where targets come from
 - idea: *define target, reduce error*

Two Core Algorithms: SARSA and Q-learning

- both learn action-value function $Q(s, a)$
 - table of size $|S| \times |A|$
- both update online from interaction
- SARSA: learn Q^π
- Q-learning: learn Q^{π^*}

SARSA & Q-learning (short)

SARSA

Repeat (for each episode):

Init s

Select $a \sim \pi(s) = \varepsilon$ -greedy on Q

Repeat (for each step of episode):

Take action a , observe r, s'

Select $a' \sim \pi(s') = \varepsilon$ -greedy on Q

$Q(s, a) += \alpha[r + \gamma Q(s', a') - Q(s, a)]$

$s \leftarrow s', a \leftarrow a'$

... until s is terminal

Q-learning

Init $Q(s, a)$ arbitrary, $Q(\text{terminal}, \cdot) = 0$

Repeat (for each episode):

Init s

Repeat (for each step of episode):

Select $a \sim \pi(s) = \varepsilon$ -greedy on Q

Take action a , observe r, s'

$Q() += \alpha[R + \gamma \max Q(s', \cdot) - Q(s, a)]$

$s \leftarrow s'$

... until s is terminal

- similar structure
- key difference: target

General scheme

Initialize $Q(s, a)$ for all $s \in S$, $a \in A(s)$ arbitrarily, $Q(\text{terminal}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):

Choose A from S using policy derived from Q (e.g. ε -greedy)

Take action A , observe R , S'

$Q(S, A) \leftarrow Q(S, A) + \alpha[G(S', A') - Q(S, A)]$

$S \leftarrow S'$

until S is terminal

- Update Q: regular delta-rule learning the target $G(s, a)$

What Is Being Learned?

- $Q(s, a)$: expected long-term return
- policy implicit or explicit
- learning via error correction
- target minus current estimate
- value = expected return estimate
- learning signal: $\delta = \text{target} - Q$
- control: make policy greedy wrt Q

Returns

Return G_t :

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

- sum of future rewards
- random variable
- high variance
- delayed signal
- depends on future randomness
- horizon grows variance
- discount: $\gamma \in [0, 1)$

State value:

$$V^\pi(s) = E_\pi[G_t \mid s_t = s] = E_\pi[G(s)]$$

Action value:

$$\begin{aligned} Q^\pi(s, a) &= E_\pi[G_t \mid s_t = s, a_t = a] \\ &= E_\pi[G(s, a)] \end{aligned}$$

- policy: $\pi(a|s)$
- expectation over trajectories
- objective: maximize $E[G_0]$

How to learn V/Q?
aka *evaluate policy*

Policy evaluation option 1: Monte Carlo Evaluation

- **idea:** estimate expected returns by sampling
- for a fixed policy π
- sample episodes
- estimate returns
- update by averaging
- learn from complete episodes
- first-visit vs every-visit
- incremental averaging view

MC update

$$Q(s, a) \leftarrow Q(s, a) + \alpha_t (G_t - Q(s, a))$$

- target = sampled return
- needs episode termination
- episode-level updates
- unbiased
- high variance

Recursive Form of Return

Return G_t :

$$\begin{aligned} G_t &= r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \\ &= r_t + \gamma [r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3}] \\ &= r_t + \gamma G_{t+1} \end{aligned}$$

- bootstrapping idea
- local target
- leads to Bellman equations
- one-step lookahead
- can be more than one -step lookahead, i.e. n-step
- bootstrap target uses estimate
- bridge to TD targets

Bellman Expectation Equation

For a fixed policy π :

$$V_{\pi}(s) = E_{\pi}[r_{t+1} + \gamma V_{\pi}(s_{t+1}) \mid s_t = s]$$

$$Q_{\pi}(s, a) = E[r_{t+1} + \gamma E_{\pi}[Q_{\pi}(s_{t+1}, a')]]$$

- maps value function to value function
- contraction
- fixed point: $Q_{\pi} = T^{\pi} Q_{\pi}$
- not an algorithm
- T^{π} contraction with factor γ
- repeated application converges
- DP uses full expectation

Policy evaluation option 2: Temporal-Difference Learning

$$\text{TD_target} = r + \gamma Q(s', a')$$

$$\text{TD_error} = \delta = \text{TD_target} - Q(s, a)$$

- sample Bellman target
- update every step
- lower variance than MC
- biased but consistent
- bootstrap: use $Q(s', a')$
- online updates per step
- lower variance targets

Practice 1: Q-learning, SARSA, MC-control

From Evaluation to Two-Stage Control scheme

- goal: optimal policy
 - improve policy using value estimates
 - alternate evaluation and improvement
 - evaluation: estimate Q_π
 - improvement: greedy wrt Q
 - generalized policy iteration idea
- stage 1: policy evaluation
 - stage 2: policy improvement
 - repeat until stable
 - evaluate approximately
 - improve greedily
 - interleave in practice

Stage 1: Fix Policy

- $Q \rightarrow$ greedy policy $\rightarrow$ epsilon-greedy
- learn Q_π
- need exploration
- behavior vs target viewpoint
- target: greedy, behavior: ε -greedy
- learn on-policy Q^π

Stage 2: Policy Improvement

- $\pi(s) = \arg \max_a Q(s, a)$
- greedy improvement
- local decision
- increases value
- improvement operator
- makes policy sharper

Convergence

- tabular case
- sufficient exploration
- decreasing step sizes
- nuances later
- sufficient exploration: visit all pairs
- step sizes: $\sum \alpha_t = \infty, \sum \alpha_t^2 < \infty$

Collapsing Stage 1

- evaluation and improvement now interleaved
- explicit policy not explicit
- greedy action selection inside loop
- policy improvement inside data collection
- evaluation never fully converged
- GPI loop: policy/value co-adapt

Back to SARSA

$$Q(s, a) \leftarrow Q(s, a) + \alpha(r + \gamma Q(s', a') - Q(s, a))$$

- on-policy
- target follows behavior
- on-policy TD control
- learns value of behavior policy
- sensitive to exploration in risky MDPs

SARSA scheme

SARSA

Initialize $Q(s, a)$ for all $s \in S$, $a \in A(s)$ arbitrarily, $Q(\text{terminal}, \cdot) = 0$

Repeat (for each episode):

Initialize S

$a \sim \pi_\epsilon(s)$: Choose a from s using policy derived from Q (e.g. ϵ -greedy)

Repeat (for each step of episode):

$s', r' = \text{env.step}(a)$: Take action a , observe r, s'

$a' \sim \pi_\epsilon(s')$: Choose a' from s' using policy derived from Q (e.g. ϵ -greedy)

$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$

$s \leftarrow s', a \leftarrow a'$

until s is terminal

Expected SARSA


Replace TD target with:

$$\begin{aligned}\text{TD_target} &= r + \gamma \mathbb{E}_{a' \sim \pi_\epsilon} [Q(s', a')] \\ &= r + \gamma \sum_{a'} \pi_\epsilon(a' | s') Q(s', a')\end{aligned}$$

- expectation instead of sample
- **lower variance**
- more computation
- bridges SARSA and Q-learning

Cliff Walking Example

- goal: reach RB cell from LB start
- dangerous states: high penalty + termination
- optimal behavior: shortest path
 - it's risky under exploration
- resulting SARSA policy depends on ϵ
 - higher $\epsilon \Rightarrow$ “safer” path



- During training exploration is turned on: π_ϵ^Q . When we turn it off to π^Q , is it optimal?
- $Q^{\text{what pi???}}(s, a)$
- Does SARSA differentiate behavior and target policies?

How could we “fix” that?

Option 1:

Set $\varepsilon \rightarrow 0$ decay during training. In limit, it converges to optimal greedy policy.

Option 2:

Instead of learning Q^{π_ε} for behavior policy, learn it for the target (=greedy) policy: Q^π .

How can we learn Q for the target greedy policy?

$$\text{exploration behavior expectation: } \text{TD_target} = r + \gamma \mathbb{E}_{a' \sim \pi_\varepsilon} Q(s', a')$$

$\Downarrow$

$$\text{greedy behavior expectation: } \text{TD_target} = r + \gamma \mathbb{E}_{a' \sim \pi} Q(s', a')$$

$$= r + \gamma Q\left(s', \arg \max_{a'} Q(s', a')\right) = r + \gamma \max_{a'} Q(s', a')$$

Bellman Optimality Equation

$$Q^*(s, a) = E[r + \gamma \max_{a'} Q^*(s', a')]$$

- optimal control objective
- optimality operator T^*
- combines evaluation + improvement
- fixed point: $Q^* = T^*Q^*$

- on-policy:
 - target = behavior
 - Bellman expectation operator
- off-policy:
 - target != behavior
 - Bellman optimality operator

Q-learning

Initialize $Q(s, a)$ for all $s \in S$, $a \in A(s)$ arbitrarily, $Q(\text{terminal}, \cdot) = 0$

Repeat (for each episode):

Initialize S

Repeat (for each step of episode):

$a \sim \pi_\epsilon(s)$: Choose a from s using policy derived from Q (e.g. ϵ -greedy)

$s', r' = \text{env.step}(a)$: Take action a , observe r, s'

$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$

$s \leftarrow s'$

until s is terminal

SARSA vs Q-learning

On-policy vs Off-policy

Q-function for



- SARSA: behavior policy
- Q-learning: greedy policy

Update

- SARSA: sample from π
- Expected SARSA: average over π
- Q-learning: maximize over actions

Different optimization force

- Bellman expectation operator doesn't optimize Q-induced policy
- SARSA: soft-greedy behavior
- Q-learning: Bellman optimality operator



Off-policy

- differentiates behavior and target policies
- enables data reuse (replay buffer)
- tabular case: strong convergence + optimality guarantees
- can be unstable with approximation (deadly triad)
 - data distribution is mismatched
 - projection norm is incorrect
 - breaks operator

Offline RL (Spoiler)

- learn from fixed data
- possible with off-policy methods
- powerful, but many pitfalls
- fixed dataset, no exploration
- distribution shift issues
- extrapolation error

Wrap-up: Control Methods

- MC control vs SARSA vs Q-learning
- same structure, different targets
- all are GPI variants
- differ by target definition
- compare: bias/variance, risk, data reuse

Bellman Operators


- expectation: $(T^\pi Q)(s, a) = E[r + \gamma Q(s', a')]$
- optimality: $(T^* Q)(s, a) = E[r + \gamma \max_{a'} Q(s', a')]$
- expectation operator
- optimality operator
- contraction mappings
- fixed points
- sampling gives TD targets

Policy Iteration and Value Iteration

- simpler setting: planning with known model
- evaluation + improvement
- PI: T^π to convergence + greedy and again
- VI: repeated T^* updates

In practice

- small tabular MDP
- full Bellman backups
- compare convergence
- synchronous vs asynchronous backups
- stop criteria: max change < tol
- compare iterations vs cost



Takeaway

- define target
- move value toward target
- control via improvement pressure
- differences are structural
- choose target: MC vs TD vs optimality
- choose coupling: on-policy vs off-policy
- one template: target, error, update